7th of Jan. 2026

Network Designs

Three - Layer Campus Design:

The three - layer campus network design is ~~hierarchical~~ hierarchical architecture commonly used to build scalable, reliable, and manageable enterprise networks. It divides the network into three functional layers: Access, Distribution and Core. Each layer has a specific role, which helps simplify network design, improve performance, and make trouble-shooting easier. By separating concerns, the model allows network engineers to optimize hardware, protocols, and policies according to the function of each layer rather than applying a one-size-fits-all ~~approach~~ approach.

The Access layer is the point where end devices such as PCs, printers, IP phones, wireless access points, and IoT devices connect to the network. Its primary function is to provide network access while enforcing contro mechanism such as port security, VLAN assignment, Quality of Service (QoS), and authentication (e.g., 802.1x). Access layer switches are typically optimized for high port density and fast Ethernet or Gigabit connectivity. This layer also marks traffic for prioritization and serves as the boundary where user traffic enters the campus network.

The Distribution layer acts as an aggregation point for multiple access layer switches and is responsible for policy enforcement and traffic control. At this layer, routing between VLANs (inter-VLAN routing) is commonly performed, along with the implementation of access control lists (ACLs), route summarization, and redundancy mechanisms such as HSRP or VRRP. The distribution layer provides fault isolation so that problems in one access segment do not impact the entire network, and it ensures efficient and controlled traffic flow toward the core.

The Core layer is the backbone of the campus network and is designed for high-speed, highly reliable data transport. Its main function is to move traffic as quickly as possible between distribution layer devices with minimal latency and without complex policy processing. Core switches are typically high-performance, redundant, and optimized for fast switching and routing. By keeping the core simple and resilient, the three-layer design ensures scalability, high availability, and predictable performance across the entire campus network.

Collapsed Core Design:

The collapsed core design is a network architecture commonly used in medium-sized campus networks where the core and distribution layers are combined into a single layer. This design is derived from the traditional three-layer campus model but simplifies it by merging the core and distribution functions into one set of high-performance switches. The primary goal of the collapsed core is to reduce cost, complexity, and management overhead while still maintaining high availability and ~~acceptablit~~ acceptable performance for the campus network.

In a collapsed core design, the combined core/distribution layer handle both high-speed traffic forwarding and policy-based control. This layer performs functions such as inter-VLAN routing, access control lists (ACLs), Quality of Service (QoS), route summarization, and redundancy protocols, while also serving as the central aggregation point for all access layer switches. Because this layer carries out critical tasks for the entire campus, it is typically built using powerful, redundant multilayer switches with high throughput, fast convergence, and multiple uplinks to the access layer to prevent single points of failure.

The access layer in a collapsed core architecture remains largely the same as in a three-layer design, providing connectivity to end devices and enforcing edge policies such as port security, VLAN assignment and traffic marking. However, since there is no seperate core layer, traffic from the access layer is forwarded directly to the collapsed core switches. This design works best when network size and traffic volume ~~does not~~ do not justify a dedicated core layer, and when simplicity, cost efficiency, and ease of management are priorities. While the collapsed core design may not scale as well as a full three-layer model, it offers an effective and practical solution form many enterprise and campus environments.

Spine-Leaf Design:

The spine-leaf design is a modern data center and campus network architecture optimized for high bandwidth, low latency, and scalability. Unlike traditional hierarchical models, spine-leaf uses a two-tier topology in which leaf switches connect directly to endpoints (servers, storage, or access switches), and spine switches act as the high-speed backbone. Every leaf switch connects to every spine switch, creating a predictable and uniform network fabric where traffic typically traverse no more than one leaf and one spine hop.

The leaf layer serves as the access point for end devices and is responsible for enforcing network policies such as VLANs, security controls, and Quality of Service (QoS). In data center environments, leaf switches often perform functions like Layer 2/ Layer 3 gateway services, server segmentation, and virtual network integration. Because leaf switches connect to all spines, they have equal-cost paths to any destination, ensuring consistent performance regardless of which endpoints are communicating.

The spine layer provides fast, resilient forwarding between leaf switches without implementing complex policies. Spine switches are designed for high ~~throught and point~~ throughput and low latency, forwarding traffic based on routing rather than switching to avoid loops and improve convergence. Technologies such as Equal-Cost Multi-Path (ECMP) routing are commonly used to load-balance traffic accross multiple spine links, maximizin bandwidth utilization and redundancy.

One of the key advantages of the spine-leaf design is its horizontal scalability. Adding more capacity is as simple as adding another leaf (to connect more devices) or another spine (to increase bandwidth across the fabric) without redesigning the network. This architecture is especially well-suited for east-west traffic patterns typical in modern data centers, cloud environments, and large-scale campus networks. By providing predictable latency, high availability, and efficient use of bandwidth, the spine-leaf design has become a foundational model for modern network infrastructures.

CamScanner